

Implementation Blueprint for a Promptable Image-Segmentation Service

Production-style reference: Dockerised REST API + CLI Client

Executive Summary

This report specifies a concrete, production-style (“Cisco-style”: clean structure, explicit governance, repeatable CI/CD, compliance-first data handling) implementation plan for a repository that hosts a promptable image-segmentation service with (a) a Dockerised REST API and (b) a small CLI client.

The emphasis is on: (1) public redistribution safety — no raw datasets, no ambiguous/proprietary weights in the repository; (2) repeatability — pinned dependencies, deterministic modes, container pinning; and (3) operational maturity — CI, releases, monitoring hooks, security policies.

Key Conservative Design Choices

- Model artefacts are **never committed to the repository**; instead they are fetched at runtime and verified via checksum. This avoids hard file limits (100 MiB blocks) and keeps the repo small and forkable.
- Model artefacts are distributed via versioned release assets by default, supporting large binaries without repo bloat.
- Datasets are consumed only via download scripts; the repo contains dataset 'cards' and attribution files, and validation scripts, but never stores raw images. This is essential because several popular segmentation datasets have licence constraints or provenance requirements.
- CI workflows are pinned, least-privilege, and supply-chain aware.

Repo Skeleton and Files

Target Repo Layout

A production-friendly layout that cleanly separates product code, MLOps scripts, infrastructure, monitoring assets, and documentation:

```
.
├── README.md
├── pyproject.toml
├── LICENSE
├── NOTICE
├── SECURITY.md
├── CONTRIBUTING.md
├── CODE_OF_CONDUCT.md
├── CHANGELOG.md
├── .gitignore
├── .dockerignore
├── .pre-commit-config.yaml
├── .editorconfig
├── Makefile
└── docker/
    ├── Dockerfile
    └── entrypoint.sh
├── configs/
    ├── service.yaml
    ├── logging.yaml
    ├── model_registry.json
    └── prometheus/
        └── prometheus.yml
└── data/
    ├── README.md
    ├── metadata/
    │   ├── datasets/
    │   │   ├── openimages_v5_masks.md
    │   │   ├── coco_2017_instances.md
    │   │   ├── voc2012_segmentation.md
    │   │   ├── davis2017.md
    │   │   ├── youtube_vos.md
    │   │   ├── model_cards/
    │   │   │   ├── sam.md
    │   │   │   ├── sam2.md
    │   │   │   └── mobilesam.md
    │   │   └── attribution/
    │   │       ├── DATA_SOURCES.md
    │   │       └── THIRD_PARTY_NOTICES.md
    ├── scripts/
    │   ├── datasets/
    │   │   ├── download_openimages_v5.sh
    │   │   ├── download_coco_2017.sh
    │   │   ├── download_voc2012.sh
    │   │   ├── download_davis2017.sh
    │   │   ├── download_youtube_vos.sh
    │   │   └── verify_checksums.py
    │   ├── export_onnx.py
    │   └── package_release_assets.sh
└── src/
    └── son_segserve/
        ├── __init__.py
        ├── api/
        │   ├── app.py
        │   ├── routes.py
        │   └── schemas.py
```

```
    cli/
      main.py
      core/
        config.py
        logging.py
        security.py
      model/
        loader.py
        registry.py
        prompts.py
        postprocess.py
      metrics/
        prometheus.py
        drift.py
      utils/
        hashing.py
        io.py
  monitoring/
    grafana/
    dashboard.json
    alerts/
    rules.yml
  tests/
    unit/
    integration/
    regression/
    synthetic/
  docs/
    index.md
    architecture.md
    api.md
    datasets.md
    models.md
    contributing.md
    runbooks.md
```

Design rationale: Keep data as metadata + scripts only — this supports public redistribution while remaining legally conservative. Keep infra 'portable': Docker + optional compose + sample Prometheus/Grafana configs.

File-to-Purpose Mapping

File / Directory	Purpose	Notes / Acceptance Criteria
README.md	Product definition, quickstart, demo, governance entry points	Must contain 'no datasets/weights in git' note
pyproject.toml	Single source of truth for packaging, dependencies, scripts	Supports reproducible local + CI installs
LICENSE	Repo's code licence	Use a standard licence (Apache-2.0 or MIT recommended)
NOTICE	Attribution + legal notices	Include third-party notices + dataset attribution pointers
SECURITY.md	Vulnerability reporting policy	Include contact + response SLA
CONTRIBUTING.md	Contributor workflow and expectations	Recognisable standard location
data/metadata/datasets/*.md	Dataset cards for transparency and compliance	Include licence, splits, checksums, provenance notes
scripts/datasets/*	Repeatable dataset download + verification	Never commit dataset binaries
configs/model_registry.json	Machine-readable registry of model artefacts and checksums	Enables download-on-first-run verification
monitoring/*	Dashboards + scrape configs + alerts stub	Grafana dashboards are JSON models

Minimal Content Templates

README.md

```
# SON SegServe – Promptable Image Segmentation Service

A production-style reference implementation of a promptable segmentation service:
- REST API (containerised)
- CLI client (thin wrapper over the API)
- Model artefacts fetched at runtime + checksum-verified (no weights committed)

## Quickstart (Docker)
```bash
docker build -f docker/Dockerfile -t son-segserve:dev .
docker run --rm -p 8080:8080 -e SON_MODEL_ID=sam2_tiny son-segserve:dev
```

## Quickstart (CLI)
```bash
pip install -e ".[dev]"
son-segserve ping --api http://localhost:8080
son-segserve segment --api http://localhost:8080 --image ./demo.jpg \
--prompt box:10,10,200,200
```

## Repo Principles
- Do not commit datasets or model weights to the repository.
- Use scripts/datasets/* to download open datasets locally.
- Use versioned release assets for model checkpoints + checksums.

## Docs
- Architecture: docs/architecture.md
- API: docs/api.md
- Datasets & licensing: docs/datasets.md
```

pypackage.toml

```
[project]
name = "son-segserve"
version = "0.1.0"
description = "Promptable image segmentation service (REST + CLI)"
requires-python = ">=3.10"
readme = "README.md"
license = {file = "LICENSE"}
authors = [{name = "SON Contributors"}]
dependencies = [
    "fastapi>=0.110",
    "uvicorn[standard]>=0.27",
    "pydantic>=2.6",
    "httpx>=0.27",
    "numpy>=1.26",
    "pillow>=10.0",
    "prometheus-client>=0.20",
    # torch / onnxruntime intentionally not pinned here:
    # prefer extra groups per backend/hardware
]

[project.optional-dependencies]
dev = [
    "pytest>=8.0",
    "pytest-cov>=5.0",
    "ruff>=0.4",
    "mypy>=1.8",
```

```
"pre-commit>=4.0",
]
cpu = [ "onnxruntime>=1.17" ]
gpu = [ "onnxruntime-gpu>=1.17" ]

[project.scripts]
son_segserve = "son_segserve.cli.main:app"
```

```
[tool.ruff]
line-length = 100
```

```
[tool.pytest.ini_options]
addopts = "-q"
testpaths = ["tests"]
```

.gitignore (public-ML-safe defaults)

```
# Python
__pycache__/
*.pyc
.venv/
.pytest_cache/
.mypy_cache/
.ruff_cache/

# Local data / checkpoints (never commit)
data/raw/
data/processed/
data/cache/
models/
*.pt
*.pth
*.onnx
*.ckpt

# OS/IDE
.DS_Store
.idea/
.vscode/

# Build outputs
dist/
build/
*.egg-info/
```

.dockerignore

```
.git
.venv
__pycache__/
*.pyc
data/raw
data/processed
models
dist
build
tests
docs
```

.pre-commit-config.yaml

```
repos:
  - repo: https://github.com/pre-commit/pre-commit-hooks
    rev: v5.0.0
```

```
hooks:
  - id: end-of-file-fixer
  - id: trailing-whitespace
  - id: check-yaml
- repo: https://github.com/astral-sh/ruff-pre-commit
  rev: v0.6.4
  hooks:
    - id: ruff
    - id: ruff-format
```

SECURITY.md

```
# Security Policy

## Supported Versions
Only the latest `main` branch and the latest tagged release are supported.

## Reporting a Vulnerability
Please do NOT open a public issue.
Email: security@example.org
Include:
- version/tag/commit SHA
- reproduction steps
- impact assessment (if known)
We aim to acknowledge reports within 72 hours.
```

CONTRIBUTING.md

```
# Contributing

## Development Setup
python -m venv .venv && source .venv/bin/activate
pip install -e ".[dev]"
pre-commit install
pytest

## Pull Request Process
- One feature/fix per PR
- Add/adjust tests
- Update docs if behaviour changes
- Ensure ruff, mypy, and pytest pass locally

## Code Style
Run `pre-commit run -a` before pushing.
```

NOTICE

SON SegServe
Copyright (c) 2026 SON Contributors

This distribution includes third-party components; see
data/attribution/THIRD_PARTY_NOTICES.md.

Dataset and model licences are documented in data/metadata
and must be reviewed before commercial use.

CI/CD Workflows

The following YAML outlines reflect three non-negotiables:

- Least-privilege workflow permissions (job-level).
- Pin Actions to commit SHAs.
- Do not upload raw datasets; only upload build artefacts and small test fixtures.

CI Workflow (lint + unit + integration)

```
name: CI
on:
  pull_request:
  push:
    branches: [main]

permissions:
  contents: read

jobs:
  test:
    runs-on: ubuntu-latest
    strategy:
      matrix:
        python: ["3.10", "3.11"]
    steps:
      - name: Checkout
        uses: actions/checkout@v5
      - name: Set up Python
        uses: actions/setup-python@v5
        with:
          python-version: ${{ matrix.python }}
      - name: Install
        run: |
          python -m pip install -U pip
          pip install -e ".[dev]"
      - name: Lint
        run: |
          ruff check .
          ruff format --check .
      - name: Type-check
        run: mypy src
      - name: Unit tests
        run: pytest -q --cov=son_segserve --cov-report=term-missing

integration:
  runs-on: ubuntu-latest
  needs: [test]
  steps:
    - uses: actions/checkout@v5
    - run: docker build -f docker/Dockerfile -t son-segserve:ci .
    - run: |
        docker run -d --name api -p 8080:8080 son-segserve:ci
        python -m pip install -U httpx
        python scripts/ci_smoke.py --base-url http://localhost:8080
```

Docker Build + Push Workflow (Container Registry)

```
name: Docker
on:
  release:
    types: [published]
```

```

env:
  REGISTRY: ghcr.io
  IMAGE_NAME: ${{ github.repository }}

jobs:
  build_push:
    runs-on: ubuntu-latest
    permissions:
      contents: read
      packages: write
      attestations: write
      id-token: write
    steps:
      - uses: actions/checkout@v5
      - name: Login to Container Registry
        uses: docker/login-action@<PINNED_SHA>
        with:
          registry: ${{ env.REGISTRY }}
          username: ${{ github.actor }}
          password: ${{ secrets.GITHUB_TOKEN }}
      - name: Extract metadata
        id: meta
        uses: docker/metadata-action@<PINNED_SHA>
        with:
          images: ${{ env.REGISTRY }}/{{ env.IMAGE_NAME }}
      - name: Build and push
        id: push
        uses: docker/build-push-action@<PINNED_SHA>
        with:
          context: .
          file: docker/Dockerfile
          push: true
          tags: ${{ steps.meta.outputs.tags }}
          labels: ${{ steps.meta.outputs.labels }}
      - name: Attest provenance
        uses: actions/attest-build-provenance@v3
        with:
          subject-name: ${{ env.REGISTRY }}/{{ env.IMAGE_NAME }}
          subject-digest: ${{ steps.push.outputs.digest }}
          push-to-registry: true

```

Note: Keep images lean via multi-stage builds and exclude caches to stay within layer limits.

Docs Deploy Workflow

```

name: Docs
on:
  push:
    branches: [main]

permissions:
  contents: read
  pages: write
  id-token: write

jobs:
  build:
    runs-on: ubuntu-latest
    steps:
      - uses: actions/checkout@v5
      - uses: actions/setup-python@v5
        with:
          python-version: "3.11"
      - run: |
          python -m pip install -U pip

```

```

    pip install -e ".[dev]"
    pip install mkdocs mkdocs-material
    mkdocs build
  - name: Upload Pages artifact
    uses: actions/upload-pages-artifact@v3
    with:
      path: site/
deploy:
  needs: build
  runs-on: ubuntu-latest
  environment:
    name: github-pages
    url: ${{ steps.deployment.outputs.page_url }}
steps:
  - name: Deploy
    id: deployment
    uses: actions/deploy-pages@v4

```

Release Workflow

```

name: Release
on:
  workflow_dispatch:
  push:
    tags: ["v*.*.*"]
permissions:
  contents: write

jobs:
  cut_release:
    runs-on: ubuntu-latest
    steps:
      - uses: actions/checkout@v5
      - name: Build model artefact bundle
        run: |
          bash scripts/package_release_assets.sh
          ls -lh dist/
      - name: Create release and upload assets
        env:
          GH_TOKEN: ${{ github.token }}
        run: |
          gh release create "${GITHUB_REF_NAME}" \
            --title "${GITHUB_REF_NAME}" \
            --notes "SON SegServe release ${GITHUB_REF_NAME}" \
            dist/*

```

Local Equivalent (Maintainer Runbook)

```

gh auth login
git tag v0.1.0 && git push origin v0.1.0
gh release create v0.1.0 dist/*

```

Model and Dataset Artefact Handling

Why You Must Not Commit Weights or Datasets

- Repositories block files over 100 MiB in normal pushes and recommend keeping repos small (ideal < 1 GB; strongly recommended < 5 GB).
- Git LFS has plan-dependent per-file limits and still encourages careful storage strategy.
- **Therefore: ship download scripts + checksums, not binaries.**

Model Registry Pattern (configs/model_registry.json)

```
{  
    "schema_version": 1,  
    "models": [  
        {  
            "model_id": "sam2_tiny",  
            "family": "sam2",  
            "backend": "onnxruntime",  
            "format": "onnx",  
            "artifact": {  
                "uri": "https://<host>/releases/download/v0.1.0/sam2_tiny.onnx",  
                "sha256": "REPLACE_WITH_SHA256",  
                "size_bytes": 123456789  
            },  
            "license": "requires_review",  
            "notes": "Downloaded on first run; cached under XDG cache dir."  
        }  
    ]  
}
```

Download-on-First-Run + Checksum Verification

Recommended cache location:

- Default: \${XDG_CACHE_HOME:-~/cache}/son-segserve/models/
- Override via: SON_CACHE_DIR=/path

```
import hashlib  
from pathlib import Path  
  
def sha256_file(path: Path, chunk_size: int = 1024 * 1024) -> str:  
    h = hashlib.sha256()  
    with path.open("rb") as f:  
        for chunk in iter(lambda: f.read(chunk_size), b""):  
            h.update(chunk)  
    return h.hexdigest()  
  
def verify_sha256(path: Path, expected_hex: str) -> None:  
    actual = sha256_file(path)  
    if actual.lower() != expected_hex.lower():  
        raise ValueError(  
            f"SHA256 mismatch for {path.name}: {actual} != {expected_hex}"  
        )  
    # Principle: fail closed – do not load mismatched artefacts.
```

Artefact Distribution Options

| Option | Pros | Cons / Limits | Best Fit |
|--------|------|---------------|----------|
|--------|------|---------------|----------|

| | | | |
|----------------------------------|--|---|---|
| Versioned Release Assets | Simple UX; integrates with tags; supports large binaries; no total bandwidth/size limits; 2 GiB per asset; up to 1000 assets per release | Asset management is per-release; must maintain checksums; 2 GiB per file cap | Default for model weights + ONNX exports |
| OCI Container Registry | Good for containerised distribution; integrates with CI publishing | Per-layer limits; requires registry auth and careful layering | Shipping the service image; optional for 'model-as-OCI' |
| External Object Storage (S3/GCS) | Scales to very large artefacts; independent lifecycle | Requires account + credentials + cost + policy; increases onboarding complexity | Optional for enterprise forks; avoid in baseline |

Dataset Handling Rules and Licence Flags

Introduce a dataset compliance gate: every dataset script must write a data/processed/<dataset>/MANIFEST.json containing:

- Download URLs
- File checksums (where feasible)
- Licence field (ok_for_demo, noncommercial, requires_per_item_verification, etc.)
- Citation guidance

Licence Caveats per Dataset

- **Open Images:** annotations are CC BY 4.0; images are listed as CC BY 2.0, but publishers disclaim warranties and advise verifying each image licence.
- **YouTube-VOS:** annotations under CC BY 4.0, but the dataset is released for non-commercial research only — generally incompatible with a broadly redistributable public demo intended for mixed use. Mark as 'requires review / likely exclude from default demo.'
- **VOC2012:** images sourced from Flickr; use must respect Flickr terms of use (licence is not uniformly permissive). Red flag for public demo assets committed to the repo.
- **DAVIS:** train+val is 90 sequences; evaluation uses region J and boundary F. Good for video-seg evaluation but not a 'tiny quickstart'.
- **COCO:** large-scale instance segmentation masks. Using subsets for demos is recommended to keep download weight manageable.

Conservative Recommendations for Demo Defaults

- **Default demo dataset:** Open Images masks (because mask annotations are explicit and large-scale) but with a prominent 'verify image licences' warning.
- **Default demo fixtures committed to repo:** synthetic images you generate yourself (no third-party licence risk).
- Keep YouTube-VOS and VOC off by default in scripts unless the user explicitly accepts their constraints.

Governance and Repo Lifecycle

Repository Lifecycle Flow

```
graph TD; dev[Developer workstation] --> pr[Pull Request]; pr --> ci[CI: lint + tests + security checks]; ci -->|pass| merge[Merge to main]; merge --> build[Build: container + docs]; build --> release[Release: tag + release assets]; release --> deploy[Deploy: VM / container runtime]; deploy --> monitor[Monitor: metrics + logs + drift]; monitor --> dev
```

Branch Protection Rules (Recommended)

- Require PR reviews before merging
- Require status checks before merging
- Require conversation resolution before merging
- Require signed commits (optional but recommended for high-integrity repos)
- Require linear history (optional; helps traceability)

Also recommended: add a CODEOWNERS file and enable 'Require review from Code Owners' to ensure sensitive areas (security, CI, release scripts) are reviewed.

Suggested Branch Naming Strategy

- main (protected)
- feature/<short-scope>
- fix/<issue-id>-<short-scope>
- release/<version> (optional, if you gate releases separately)

Recommended Security Settings

- Secret scanning — enable for public repositories.
- Dependency update automation via dependabot.yml.
- Code scanning (CodeQL default setup) — eligible for public repos with CI enabled; scans on pushes/PRs and on a schedule.

Tests, Monitoring, Security, Reproducibility

Test Suite Inventory

Keep tests small, deterministic, and licence-safe:

- **Unit tests:** pure functions (hashing, config parsing, prompt parsing, mask post-process).
- **Integration tests:** start container, hit /healthz, /segment, /metrics.
- **Model regression ('golden set'):** use generated synthetic images + fixed prompts; store expected summary statistics (mask area, connected-components count) instead of raw masks to avoid brittle exact pixel equality. Track a hash of the model artefact + inference backend version in the regression output.
- **Data validation tests:** validate dataset manifests and expected directory layouts after download (do not run full dataset downloads in CI).
- **Synthetic tests:** fuzz prompts (invalid box coords, empty points, huge images) to harden request validation.

Local Run Commands (Developer Runbook)

```
pip install -e ".[dev]"
pre-commit run -a
pytest
docker build -f docker/Dockerfile -t son-segserve:dev .
docker run --rm -p 8080:8080 son-segserve:dev
```

Monitoring and Observability Hooks

Metrics Endpoint Spec

Expose GET /metrics in Prometheus text format with Content-Type: text/plain; version=0.0.4.

Metrics to collect (minimum viable):

- HTTP request count + status
- Request latency histogram (p50/p95 derived in Prometheus/Grafana)
- Inference time histogram
- Memory RSS gauge (process metric)
- Mask count per request (gauge/hist)
- Predicted quality proxy (e.g. predicted_iou distribution)

Prometheus Scrape Config

```
global:
  scrape_interval: 15s

scrape_configs:
  - job_name: son-segserve
    metrics_path: /metrics
    static_configs:
      - targets: ["host.docker.internal:8080"]
```

Sample Grafana Dashboard JSON Stub

```
{
  "id": null,
  "uid": "son-segserve",
  "title": "SON SegServe",
  "tags": ["segmentation", "son"],
  "timezone": "browser",
```

```

"schemaVersion": 41,
"panels": [
  {
    "type": "timeseries",
    "title": "HTTP latency (p95)",
    "targets": [
      {
        "expr": "histogram_quantile(0.95, sum(rate(http_request_duration_seconds_bucket[5m])) by (le))"
      }
    ],
    "gridPos": { "h": 8, "w": 24, "x": 0, "y": 0 }
  }
]
}

```

Security and Licensing Checks

Licence Auditing Workflow

- 1) **Model licence:** verify upstream terms (code + checkpoints) and write a model card under data/metadata/model_cards/. If any ambiguity exists, set license: requires_review and do not include weights in default releases.
- 2) **Dataset licence:** record dataset homepage, licence text, and restrictions in dataset card.
- 3) **Repo licence:** choose a clear repo licence (Apache-2.0 or MIT). Licence detection works with standard filenames.
- 4) **Enforce a 'no large files' policy:** CI check that fails if a PR adds *.pt, *.onnx, or files above a threshold.

Reproducibility Steps

- Record seeds and deterministic settings in configs/service.yaml and print them at startup.
- Add a runtime --deterministic flag: enables deterministic algorithms where possible and documents potential speed impacts.
- Pin container base images and use multi-stage Dockerfiles as per Docker best practices.
- Pin CI Actions to commit SHAs.
- For releases: include artefact checksums alongside binaries as separate *.sha256 assets.

Dockerfile (Recommended Baseline)

Multi-stage builds and `.dockerignore` are first-class Docker features that reduce image size.

```
# syntax=docker/dockerfile:1
FROM python:3.11-slim AS build
WORKDIR /app

RUN python -m pip install -U pip
COPY pyproject.toml README.md LICENSE /app/
COPY src/ /app/src/
RUN pip install --no-cache-dir -e ".[cpu]"

FROM python:3.11-slim AS runtime
WORKDIR /app

# Non-root user is recommended for runtime images
RUN useradd -m -u 10001 appuser
USER appuser

COPY --from=build /usr/local /usr/local
COPY --from=build /app/src /app/src
COPY configs/ /app/configs
COPY docker/entrypoint.sh /app/entrypoint.sh

EXPOSE 8080
ENV SON_CONFIG=/app/configs/service.yaml
ENTRYPOINT [ "/app/entrypoint.sh" ]
```

docker/entrypoint.sh

```
#!/usr/bin/env bash
set -euo pipefail
exec python -m son_segserve.api.app --host 0.0.0.0 --port 8080
```

REST API Specification (Minimum Viable)

Endpoints (versioned)

- GET /healthz → service health
- GET /readyz → readiness (model loaded, cache ok)
- GET /metrics → Prometheus metrics (text format)
- POST /v1/segment → run segmentation from image + prompt
- POST /v1/segment/video → optional: short video segmentation with temporal smoothing (advanced)

POST /v1/segment — Request JSON

```
{  
  "image": {  
    "content_type": "image/jpeg",  
    "base64": "..."  
  },  
  "prompt": {  
    "type": "box",  
    "box_xyxy": [10, 20, 200, 220]  
  },  
  "options": {  
    "return_format": "rle",  
    "max_masks": 3,  
    "stabilize": false  
  }  
}
```

Response JSON

```
{  
  "model_id": "sam2_tiny",  
  "masks": [  
    {  
      "mask_id": "0",  
      "encoding": "rle",  
      "rle": "...",  
      "predicted_iou": 0.87,  
      "area_px": 12345  
    }  
  ],  
  "timing_ms": {  
    "preprocess": 4.1,  
    "inference": 31.7,  
    "postprocess": 7.9,  
    "total": 43.7  
  }  
}
```

CLI Commands (Thin Client)

```
son-segserve ping --api http://localhost:8080  
  
son-segserve segment \  
  --api http://localhost:8080 \  
  --image ./x.jpg \  
  --prompt "box:10,20,200,220" \  
  --out ./mask.json  
  
son-segserve bench \  

```

```
--api http://localhost:8080 \
--images ./fixtures \
--prompt "point:120,88,fg"
```

PR Checklist

| Category | Checklist Item | How Reviewer Verifies |
|-----------------|--|-----------------------------------|
| Scope | PR is single-purpose and titled clearly | Title + diff inspection |
| Tests | New/changed behaviour has tests | pytest passes locally/CI |
| CI hygiene | No new workflow secrets added without justification | .github/workflows diff |
| Artefact safety | No weights/datasets/binaries committed | Repo diff + CI 'large file guard' |
| Licensing | Dataset/model licence implications documented (if touched) | Dataset/model card updated |
| Reproducibility | Seeds/config updated when changing model behaviour | configs/service.yaml + logs |
| Observability | Metrics/logging updated for new paths | /metrics samples + logs |
| Docs | README/docs updated if UX changes | Docs diff + preview build |

Assumptions

- Cloud provider is unspecified → recommend 'portable baseline' with a single container on a generic VM (or local).
- Hardware is unspecified → assume GPU optional; provide CPU ONNXRuntime path by default, GPU as an extra.
- Exact model size is unspecified → assume multiple model IDs (tiny/base/large) via model_registry.json.
- Licence clarity for some model checkpoints is unspecified → model licences are marked requires_review unless verified by an official artefact you explicitly approve.
- Public repo requirement implies: do not commit training data, do not commit weights unless you have explicit redistribution rights, and do not include non-commercial-only datasets in the default demo path.

Primary Sources and Dataset References

Model Primary Links

- SAM: <https://segment-anything.com/>
- SAM 2: <https://ai.meta.com/sam2/>
- MobileSAM: <https://arxiv.org/abs/2306.14289>

Dataset Primary Links

- COCO: <https://cocodataset.org/>
- PASCAL VOC: <http://host.robots.ox.ac.uk/pascal/VOC/>
- Open Images: <https://storage.googleapis.com/openimages/web/index.html>
- DAVIS: <https://davischallenge.org/>
- YouTube-VOS: <https://youtube-vos.org/>

Technical Documentation

- Prometheus configuration: <https://prometheus.io/docs/prometheus/latest/configuration/configuration/>
- Prometheus exposition formats: https://prometheus.io/docs/instrumenting/exposition_formats/
- Grafana dashboard JSON model: <https://grafana.com/docs/grafana/latest/reference/dashboard/>
- Docker best practices: https://docs.docker.com/develop/develop-images/dockerfile_best-practices/
- pre-commit: <https://pre-commit.com/>